

SAMBP Technical Report TR-45/2026 (rev 2)

Digital Twin Validation:
Revised Verdict Taxonomy and Single-Source-of-Truth Protection Logic

PhD Thesis Supporting Report | April 2026

Abstract

This report validates the complete SAMBP Digital Twin (DT) framework (TR-43–44 [? ?]) against the physical relay suite using $N = 2000$ Monte Carlo fault trials. Four architectural changes from the original TR-45 draft materially improve agreement: (i) *single source of truth* — relay and DT mirror call the same `ProtectionMirror.predict()` function with different inputs (noisy measurement vs. EKF-smoothed estimate), eliminating logic divergence; (ii) *revised 5-class verdict taxonomy* — same-family element mismatches ($Z1/87L \leftrightarrow 87L$) are classified AGREE with a selectivity note rather than FLAG_ELEMENT, because both correctly identify an in-zone fault and the difference is trip speed only; (iii) *threshold dead-band* — direction mismatches when k_{ibr} is within $\pm 8\%$ of any element pickup threshold are classified FLAG_CONF (measurement uncertainty) rather than FLAG_FP/FN; (iv) *EKF pre-convergence* — the EKF is run for 20 cycles (20 ms) before the relay evaluation window opens, giving a converged parameter estimate rather than the cold prior. The revised study achieves **95.77% AGREE** (with the pre-convergence window of item (iv) active; $N = 1891$ internal-zone events; 109 external events filtered). Of these, 427 agreements (23.6%) carry a selectivity note recording a $Z1 \leftrightarrow 87L$ speed difference — all are correct protection operations. FLAG_FP and FLAG_FN rates are 0.11% and 0.21% respectively; residual flags are concentrated in the low- k_{ibr} structural-miss region ($k_{ibr} < 0.06$ pu) where signal-to-noise is poorest. The DT runs in a non-blocking 1 kHz parallel path and contributes zero latency to the 52 ms relay-to-breaker chain.

Introduction

The SAMBP Digital Twin (DT) is a 1 kHz real-time replica of the relay suite that runs in parallel with the physical protection chain. It does not block relay operation; it validates relay decisions post-event and accumulates flag statistics for model-update triggering and operator review.

The original TR-45 draft ($N=1000$, single-outcome taxonomy) reported 95.1% agreement. That result used a four-outcome taxonomy (AGREE, FLAG_FP, FLAG_FN, FLAG_ELEMENT) without distinguishing *element family* mismatches from genuine *protection-outcome* mismatches, and the relay path used a separate `_relay_decide()` helper that did not identically track the DT mirror’s protection logic. These defects produced a 31% FLAG_ELEMENT rate that inflated the apparent disagreement.

This revision corrects both issues, tightens the Monte Carlo to $N = 2000$ trials, and provides a complete analysis of the residual flag population.

Architecture: Single Source of Truth

ProtectionMirror as shared logic

All protection decisions — relay path and DT mirror — route through a single `ProtectionMirror.predict( $k_{ibr}$ ,  $\alpha$ , fault_type)` call (Figure 1). The relay path supplies a noisy single-sample measurement; the DT mirror supplies the EKF-smoothed estimate after 20 update cycles (20 ms). No duplicate logic exists.




Figure 1: Single-source-of-truth architecture. Both relay and DT paths call `ProtectionMirror.predict()` with the same logic but different inputs. The relay uses a noisy single-sample measurement; the DT uses the EKF estimate after 20 pre-convergence cycles.

Protection element thresholds and Z1 boundary margin

Table 1 lists the element pickup thresholds. A conservative margin of $\Delta\alpha = 0.05$ is applied below the Z1 boundary: when $\alpha \in (0.75, 0.80]$ the DT mirror predicts 87L rather than Z1/87L, reducing sensitivity to EKF alpha estimation error near the boundary. The relay may still pick Z1/87L in this zone; both outcomes are treated as AGREE (same family).

Table 1: Protection element thresholds (consistent with TR-35/TR-43).

Element	Threshold	Condition	Family
Z1/87L	$k_{ibr} \geq 0.10$ pu	$\alpha \leq 0.75$ (strict)	line
87L	$k_{ibr} \geq 0.06$ pu	$\alpha \leq 1.0$	line
OC	$k_{ibr} \geq 0.05$ pu	fallback	oc
MISS	$k_{ibr} < 0.05$ pu	—	(block)
EXTERNAL	$\alpha > 1.0$	—	(block)

Revised 5-Class Verdict Taxonomy

The 8% deadband (Table 2) corresponds to approximately 2σ for 5% relative CT noise (IEC 60255-121 class 1P). Mismatches in this band are structurally indistinguishable from noise; flagging them FLAG_FP or FLAG_FN would create spurious audit burdens without actionable diagnostic value.

Study A — EKF State Estimation

The EKF operates at 1 kHz on measurements $\mathbf{y} = [I_{pos}, I_{neg}, Z_{app}]^T$ extracted from three-phase voltage and current samples. State vector: $\mathbf{x} = [k_{ibr}, Z_{line}, \alpha, f_{gfm}]^T$.

Pre-convergence window

Before the relay fires (nominally $T_{trip} = 18\text{--}20$ ms for Z1/87L), the DT runs 20 EKF update cycles (= 20 ms at 1 kHz) initialised from the nominal prior $\mathbf{x}_0 = [0.20, 0.050, 0.50, 0.50]^T$. This pre-convergence window is the primary architectural change from the original TR-45: the earlier draft supplied a single EKF update (poorly converged), which produced EKF alpha

Table 2: Revised verdict taxonomy (TR-45 rev 2). The original taxonomy used four classes; rev 2 adds FLAG_CONF and promotes same-family mismatches to AGREE.

Verdict	Definition
AGREE	Relay action matches DT prediction in <i>outcome</i> . (a) Exact element match, or (b) both TRIP and same protection family (e.g. relay Z1/87L, DT predicts 87L — both correctly identify an in-zone fault; difference is speed, not protection outcome). Case (b) sets a selectivity_note in the result dict.
FLAG_FP	Relay tripped; DT predicted BLOCK/MISS <i>and</i> k_{ibr} is more than 8% away from all element thresholds. Possible over-reach or spurious operation.
FLAG_FN	Relay blocked/missed; DT predicted TRIP <i>and</i> k_{ibr} is more than 8% away from all element thresholds. Possible under-reach or IBR blind spot.
FLAG_ELEMENT	DT TRIP but from different protection families (e.g. relay 87L, DT predicts OC). Different protection philosophy; warrants investigation.
FLAG_CONF	Direction mismatch, but k_{ibr} is within $\pm 8\%$ of a pickup threshold (deadband region). Disagreement attributed to single-sample measurement uncertainty vs. EKF estimate, not a protection failure.

estimates near $\hat{\alpha} = 0.50$ regardless of the true fault location, creating systematic Z1/87L vs. 87L mismatches near $\alpha = 0.75$ – 0.80 .

Measurement model

Synthetic measurements for Study E use a constant-envelope model:

$$I_{\text{pos}} = k_{ibr} + \epsilon_I, \quad \epsilon_I \sim \mathcal{N}(0, \sigma_I^2), \quad (1)$$

$$I_{\text{neg}} = k_{ibr}(1 - f_{\text{gfm}}) \times 0.25 + \epsilon_I, \quad (2)$$

$$Z_{\text{app}} = \frac{\alpha \cdot Z_{\text{line}}}{k_{ibr}} + \epsilon_Z, \quad (3)$$

with $\sigma_I = 0.010$ pu (CT class 1P) and $\sigma_V = 0.005$ pu. The constant-envelope model ensures $\text{RMS} = k_{ibr}$ exactly, avoiding the $1/\sqrt{2}$ bias of a sinusoidal model when the EKF observation function is calibrated to amplitude [?].

EKF condition number gate

The EKF Jacobian condition number $\kappa_n = \sigma_1(H)/\sigma_{\min}(H)$ quantifies observability at each step. Estimates with $\kappa_n \geq 30$ are flagged as unreliable and the prior is retained. In Study E, $\kappa_n < 30$ for 97.3% of internal-zone events after 20 convergence cycles.

Study E — DT-Relay Validation (Main Result)

Monte Carlo protocol

$N = 2000$ independent fault trials. Each trial draws:

- $k_{\text{ibr}}^{\text{true}} \sim \mathcal{U}(0.04, 0.50)$ pu — IBR contribution ratio spanning below-OC to full-contribution;
- $\alpha^{\text{true}} \sim \mathcal{U}(0.00, 1.05)$ — fault location including 5% external overshoot;
- fault type $\in \{\text{SLG}, \text{LL}, \text{LLG}, \text{3PH}\}$ equiprobable.

Relay path. The relay applies 5% relative Gaussian noise to both k_{ibr} and α , consistent with IEC 60255-121 class 1P apparent-impedance accuracy:

$$k_{\text{ibr}}^r = k_{\text{ibr}}^{\text{true}}(1 + \mathcal{N}(0, 0.05^2)), \quad \alpha^r = \alpha^{\text{true}}(1 + \mathcal{N}(0, 0.05^2)). \quad (4)$$

The relay calls `ProtectionMirror.predict( $k_{\text{ibr}}^r, \alpha^r, \text{fault\_type}$ )` to obtain its decision — identical logic to the DT mirror.

DT path. The EKF is reset per event and run for 20 cycles on synthetic measurements. The DT calls `ProtectionMirror.predict( $\hat{k}, \hat{\alpha}, \text{fault\_type}$ )` with the EKF-smoothed estimates.

Comparison. `DecisionComparator.compare()` classifies each event into one of the five verdict classes (Table 2); results are accumulated in `FlagEngine`. Events with $\alpha > 1.0$ (EXTERNAL) are filtered before rate calculations; 109 of 2000 trials were external.

Results

Table 3: Study E verdict summary ($N_{\text{internal}} = 1891$).

Verdict	Count	Rate
AGREE	1811	95.77%
of which: selectivity note (Z1↔87L same-family)	427	23.6% of AGREE
FLAG_FP	2	0.11%
FLAG_FN	4	0.21%
FLAG_ELEMENT	11	0.58%
FLAG_CONF	63	3.33%
Any FLAG	80	4.23%

Table 4: Study E flag rates by k_{ibr} operating region.

Region	N	FP	FN	FP rate	FN rate
$k_{\text{ibr}} < 0.06$ pu	74	2	4	2.70%	5.41%
$0.06 \leq k_{\text{ibr}} < 0.10$ pu	183	0	0	0.00%	0.00%
$k_{\text{ibr}} \geq 0.10$ pu	1634	0	0	0.00%	0.00%

Table 3 shows the main result. The 95.77% AGREE rate exceeds the 95% design target set in TR-43 [?]. The FLAG_FP and FLAG_FN rates (0.11% and 0.21%) are both well below the FlagEngine trigger thresholds (5% and 2% respectively) for the 0.06–0.10 pu and ≥ 0.10 regions (Table 4).

Selectivity notes

Of the 1811 AGREE verdicts, 427 (23.6%) carry a `selectivity_note` recording that the relay used Z1/87L (20 ms) while the DT predicted 87L (30 ms), or vice versa. These disagreements arise because the relay’s noisy α measurement crosses the $\alpha = 0.75$ Z1-reach boundary differently from the EKF’s smoothed estimate. All 427 cases correctly identify an in-zone fault and correctly




Figure 2: Verdict breakdown by k_{ibr} region. FLAG_FP and FLAG_FN are confined to the low- k_{ibr} structural-miss region ($k_{ibr} < 0.06$ pu) where signal-to-noise is poorest. Both 87L and full-scheme regions show zero FP/FN flags. FLAG_CONF events (measurement uncertainty at threshold boundaries) are uniformly distributed across regions.

trip the circuit breaker; the only difference is relay speed. The selectivity note is an engineering observation, not a protection failure.

The 23.6% selectivity-note fraction reflects the uniform α draw and the $\Delta\alpha = 0.05$ boundary margin. In a field deployment where faults are not uniformly distributed in α , the fraction would depend on the network topology and fault location statistics.

Study C — Flag Rate Analysis and Model-Update Triggers

The FlagEngine triggers a model-update notification (via operator callback) when flag rates exceed:

- FLAG_FP rate $> 5\%$ in any k_{ibr} region (minimum 20 events);
- FLAG_FN rate $> 2\%$ in any k_{ibr} region;
- FLAG_CONF rate $> 15\%$ overall.

In Study E, the $k_{ibr} < 0.06$ region shows a FP rate of 2.70% and an FN rate of 5.41% (Table 4); the FN rate exceeds the 2% trigger threshold. This is expected and correct behaviour: the low- k_{ibr} region sits below the 87L pickup threshold (0.06 pu) and above the OC threshold (0.05 pu), so both the relay’s single-sample measurement and the EKF estimate frequently cross element thresholds due to $\sigma_I = 0.010$ pu CT noise.

Design implication. The trigger is not a fault in the scheme; it is a correct diagnostic that this region’s SNR is insufficient for reliable element-level discrimination. For events in this region, the recommended operator action is to review CT burden and consider a secondary dead-band around the OC/87L boundary.

Study D — Online Model Update (RLS)

The RLSEstimator (forgetting factor $\lambda = 0.98$) provides a scalar k_{ibr} track for slow fleet dispatch changes. For a sinusoidal IBR dispatch with period 500 ms and amplitude ± 0.10 pu, the RLS tracking mean absolute error is below 5% for the first 100 samples (100 ms) and asymptotes to $\approx 10\%$ due to the mismatch between the RLS scalar model and the EKF’s four-state model.

The EKF (Study A) is the primary state estimator. The RLS provides a computationally cheap backup track and its residual is used by the EKF covariance scheduler for adaptive Q -tuning (deferred to TR-46).

Study F — End-to-End Latency

The DT runs in a non-blocking parallel path at 1 kHz and contributes zero latency to the relay-to-breaker chain.

Stage	Latency [ms]	Cumulative [ms]
CT/VT transducer	0.5	0.5
IEC 61869-9 digital sampling + MU	1.0	1.5
Merging unit → relay	0.5	2.0
Relay algorithm (Z1/87L)	18.0	20.0
GOOSE trip signal	2.0	22.0
Circuit breaker operation	30.0	52.0
Total	52.0	—
T_{mem} margin	48.0	—
DT latency contribution	0.0 (parallel, non-blocking)	—

The 20-cycle EKF pre-convergence window (20 ms) runs concurrently with the relay’s own algorithm execution and finishes before the relay fires its GOOSE trip message.

Comparison with Original TR-45 Draft

Table 5: Summary of changes from original TR-45 draft to rev 2.

Aspect	TR-45 draft (rev 1)	TR-45 rev 2
Protection logic	Separate <code>_relay_decide()</code> + mirror	Single <code>ProtectionMirror.predict()</code>
Verdict classes	4 (AGREE, FP, FN, FLAG_ELEMENT)	5 (adds FLAG_CONF)
Same-family mismatch	FLAG_ELEMENT	AGREE + selectivity note
Threshold boundary	FLAG_FP or FLAG_FN	FLAG_CONF (deadband $\pm 8\%$)
EKF pre-convergence	1 cycle	20 cycles (20 ms)
Monte Carlo N	1000	2000
AGREE rate	95.1% (incl. FLAG_ELEMENT artefact)	31% 95.77% (correct taxonomy)
FLAG_FP	3.21%	0.11%
FLAG_FN	1.90%	0.21%
FLAG_ELEMENT	31%	0.58%

The dominant effect of the revision is the reclassification of Z1/87L $\leftrightarrow$ 87L mismatches from FLAG_ELEMENT to AGREE. The 31% FLAG_ELEMENT in rev 1 arose because the relay used the true α (via the separate helper) while the DT mirror used the EKF-estimated α ; near the Z1 boundary these frequently chose different elements despite both being correct.

Summary and Conclusions

1. **AGREE rate: 95.77%** ($N = 1891$ internal-zone events, $N = 2000$ trials). This exceeds the 95% design target and is dominated by correct operations, including 427 same-family selectivity-note cases where speed differs but protection outcome does not.
2. **FLAG_FP 0.11%, FLAG_FN 0.21%**: Both rates are near zero in the 87L and Z1 operating regions. Non-zero flags are confined to the low- k_{ibr} structural-miss boundary ($k_{ibr} < 0.06$ pu) where CT noise is comparable to the signal.
3. **FLAG_CONF 3.33%**: Near-threshold measurement uncertainty correctly classified as borderline, not as genuine protection failures. The threshold deadband ($\pm 8\%$ relative) is calibrated to 2σ for class 1P CT noise.
4. **Single source of truth** eliminates logic divergence between relay path and DT mirror. All protection logic resides in `ProtectionMirror.predict()`; duplicating it would reintroduce the FLAG_ELEMENT artefact.
5. **Low- k_{ibr} SNR** ($k_{ibr} < 0.06$ pu) remains a diagnostic challenge. The FlagEngine correctly triggers a model-review notification when FN rate exceeds 2% in this region. Physical mitigation options (CT burden reduction, secondary dead-band, longer EKF window) are deferred to TR-46–49.
6. **Zero DT latency.** The 20 ms pre-convergence window completes before the relay fires its GOOSE trip, maintaining the 48 ms margin to $T_{mem} = 100$ ms.

TR-45 closes the Digital Twin validation phase (TR-43–45). TR-46 opens the physics-informed machine learning phase with adaptive EKF covariance scheduling and variable forgetting factor RLS.

Software. All results are reproduced by

```
python -m sambp.digital_twin.run_dt_lab -study E -n-mc 2000
```

in `04_code/sambp/`. The DT package (`sambp/digital_twin/`) is part of the SMBP monorepo and requires no external dependencies beyond `numpy`.