TSA Meeting — Slide 8 Study Note

Ch 3: SAMBP-87L — Validated Simulation Results

Contributions C1 Results · Chapter 3 · SAMBP Framework

Anoop V Eluvathingal · IIT Madras · PhD Thesis 2026

Slide Overview and Narrative

What This Slide Shows

Slide 8 is the companion results slide to Slide 7 (theory). It presents the **validated simulation results** for the complete SAMBP-87L relay across a systematic 110-scenario test matrix. The slide is organised in two panels:

Left panel — Test matrix: Four metrics compared between the fixed-slope relay and SAMBP-87L across a 4-bus test network with $k_{\text{ibr}} \in \{0.0, 0.1, 0.2, \ldots, 0.9, 1.0\}$ (11 penetration levels), three fault types (AG, BCG, ABC), and CT error up to 5%.

Right panel — Key results at a glance:

- **TPR = 1.000**: no missed trips across all 110 fault scenarios.
- **FPR = 0.000**: zero false trips under all load + CT error scenarios.
- **Robust up to** $k_{\text{ibr}} = 1.0$: fully IBR grid is covered.
- **Stage-2 gate adds only 1–2 cycle** additional decision time.

Observation: SAMBP-87L achieves perfect discrimination (TPR = 1.000, FPR = 0.000) where fixed-slope relays produce a 14.8% false-trip rate.

Full Results Table

Metric	**Fixed relay**	**SAMBP-87L**	**Gain**	**Why**
TPR (all faults)	0.783	**1.000**	+27.7%	Adaptive slope closes high-k_{ibr} gap
FPR (load / CT err)	0.148	**0.000**	−100%	Stage-2 gate vetoes uncertain trips
Avg trip time	22 ms	21 ms	−1 ms	Adaptive slope trips faster (lower k_m
k_{ibr} RMSE	—	0.024 pu	—	LM estimator tracks IBR level accur

Understanding the 110-Scenario Test Matrix

How the 110 Scenarios are Structured

The test matrix is a systematic sweep across three dimensions:

Dimension	**Values**	**Count**	**Purpose**
k_{ibr} levels	$\{0.0, 0.1, 0.2, \ldots, 1.0\}$	11	Full IBR penetration range
Fault types	AG (SLG), BCG (LL-G), ABC (3PH)	3	All asymmetric + symmetric
Internal faults	Near-end and remote-end	2	Location sensitivity
Non-fault tests	Load steps + CT error conditions	—	Security (FPR) evaluation

Fault scenarios: $11 \times 3 \times 2 = 66$ internal fault scenarios. These measure **TPR**

(true positive rate — does the relay trip when it should?).
Non-fault scenarios: 44 load + CT error conditions (balanced load steps of 0.5–3.0 pu through-current with CT error from 1–5%). These measure **FPR** (false positive rate — does the relay falsely trip when it should not?).
Total: $66 + 44 = 110$ **scenarios.**

Why $k_{\text{ibr}} = 0.0$ is included: This represents the pure synchronous generator network with no IBR. The SAMBP-87L must perform at least as well as the conventional relay in this regime. Including $k_{\text{ibr}} = 0$ confirms backward compatibility.
Why $k_{\text{ibr}} = 1.0$ is included: This represents a fully IBR grid — the asymptotic extreme. At $k_{\text{ibr}} = 1.0$, the grid-side fault current equals the IBR fault current, minimising the operate-to-bias ratio. This is the hardest possible test case. SAMBP-87L passes it with the adaptive slope reaching $k_m = k_{\text{floor}} = 0.05$.

The 4-Bus Test Network

The test network models a representative sub-transmission system:

- Generator bus (Bus A): SG at full output + IBR at rated k_{ibr}.
- Protected line: $\ell = 10$ km, $Z_L = 0.35 + j0.40\ \Omega$/km, charging $C = 0.3\ \mu$F/km.
- Remote bus (Bus B): passive load + IBR infeed.
- Load bus (Bus C): variable load for FPR testing.
- CT class: IEC Class 1P ($\epsilon_{\text{CT}} \leq 1\%$ nominal; stress-tested to 5%).

True Positive Rate: From 0.783 to 1.000

Why the Fixed Relay Misses 21.7 Percent of Faults

The fixed relay (slope $k_m = 0.30$) fails when the operate-to-bias ratio $\rho_{\min} = 2(1 - k_{\text{ibr}})/(1 + k_{\text{ibr}})$ falls below 0.30:

$$\frac{2(1 - k_{\text{ibr}})}{1 + k_{\text{ibr}}} < 0.30 \implies k_{\text{ibr}} > 0.54$$

For $k_{\text{ibr}} \in \{0.6, 0.7, 0.8, 0.9, 1.0\}$ (5 of 11 levels), the fixed relay *structurally cannot trip* for worst-case fault angle. With 3 fault types and 2 locations: $5 \times 3 \times 2 = 30$ missed scenarios. But since not all faults are worst-case angle, the actual miss count is lower: $0.217 \times 66 \approx 14$ missed scenarios, consistent with the moderate-to-high k_{ibr} range where the relay is marginal.
How SAMBP-87L recovers all 14 missed scenarios:

1. For faults at $k_{\text{ibr}} > 0.54$: the adaptive slope drops to $k_m < 0.30$, restoring the operate-to-bias margin.
2. For symmetric 3PH faults where 87L still marginal: TDCS-87L triggers on the travelling-wave transient in the first 5 ms, providing an independent trip path.
3. For asymmetric faults where k_{ibr} is very low: the 87LN negative-sequence element (OBR = 2) provides guaranteed operation.

What TPR = 1.000 Statistically Means

On 66 fault scenarios: TPR = 66/66 = 1.000. This is a point estimate. The TR-26 Monte Carlo validation (5 000 trials per group) provides the **95% Wilson confidence interval lower bound**: $\mathrm{CI}_{\mathrm{lo}} = 0.9992$, far exceeding the IEC 60255 target of ≥ 0.990. The Wilson CI formula:

$$\mathrm{CI}_{\mathrm{lo}} = \frac{\hat{p} + z^2/(2n) - z\sqrt{\hat{p}(1-\hat{p})/n + z^2/(4n^2)}}{1 + z^2/n} \tag{1}$$

with $\hat{p} = 1.000$, $z = 1.96$ (95%), $n = 5\,000$: the formula gives $\mathrm{CI}_{\mathrm{lo}} = 0.9992$. This means: with 95% statistical confidence, the true TPR is above 99.92%.

False Positive Rate: From 0.148 to 0.000

Why the Fixed Relay Produces 14.8 Percent False Trips

A "false trip" on the 87L relay occurs when the relay asserts during a non-fault condition:

- **CT error during heavy through-load**: The CT composite error generates a spurious differential current $I_{\mathrm{diff}}^{\mathrm{CT}} = 2\epsilon_{\mathrm{CT}} \cdot I_{\mathrm{bias}}$. For $I_{\mathrm{bias}} = 3.0$ pu load and $\epsilon_{\mathrm{CT}} = 5\%$ (stress test): $I_{\mathrm{diff}}^{\mathrm{CT}} = 0.30$ pu. With $I_0 = 0.10$ pu minimum pickup and fixed $k_m = 0.30$: criterion $= 0.30 \times 3.0 + 0.10 = 1.00$ pu. So at heavy load: $I_{\mathrm{diff}}^{\mathrm{CT}} = 0.30 < 1.00$ — no false trip. But at moderate IBR penetration $k_{\mathrm{ibr}} = 0.3$–0.5, where the charging current adds to the apparent differential, the threshold can be reached.
- **Charging current during energisation**: Line charging current $I_c = \omega C Z_L I_{\mathrm{bias}}/2$ appears as apparent differential. For $C = 0.3$ μF/km over 10 km at $k_{\mathrm{ibr}} = 0.4$: $I_c \approx 0.04$ pu (above minimum pickup $I_0 = 0.02$ pu). Combined with CT error: the fixed relay can false-trip during line energisation.

How SAMBP-87L achieves FPR = 0.000:

- The adaptive slope $k_m(k_{\mathrm{ibr}})$ rises at low-to-moderate k_{ibr} (where CT error is the FPR driver), providing more stability margin than the fixed relay.
- The Stage-2 veto gate catches any remaining marginal cases: CT saturation causes high κ_n (poor Jacobian condition), which drops γ below 0.85, blocking the trip.
- The combined system achieves $0/44 = 0.000$ false trips across all load and CT error scenarios in the test matrix.

The Fixed-Relay FPR Breakdown by Scenario Type

Of the $0.148 \times 44 \approx 7$ false trips by the fixed relay:

- $\sim$ 4 at high-load ($I_{\mathrm{bias}} > 2.0$ pu) with CT error $> 3\%$: charging + CT error combination.
- $\sim$ 2 during IBR reactive current injection events (not a fault, but IBR injects reactive, distorting I_{diff}).
- $\sim$ 1 at energisation with pre-existing asymmetry.

All 7 would have been caught by SAMBP-87L's adaptive slope and Stage-2 gate.

k-ibr RMSE = 0.024 pu: Estimator Accuracy

What RMSE = 0.024 pu Means for Relay Performance

The LM estimator tracks the real-time IBR penetration k_{ibr} from the received remote phasor:

$$\hat{k_{\text{ibr}}}(t) = (1-\lambda)\hat{k_{\text{ibr}}}(t-1) + \lambda \cdot \frac{|\hat{I}_R(t)|}{I_{\text{rated}}}$$

The RMSE across all 110 scenarios is:

$$\text{RMSE} = \sqrt{\frac{1}{N}\sum_{i=1}^{N}(k_{\text{ibr}}^{(i)} - \hat{k_{\text{ibr}}}^{(i)})^2} = 0.024 \text{ pu}$$

What 0.024 pu means for slope accuracy: The adaptive slope $k_m(k_{\text{ibr}}) = 0.5(1-k_{\text{ibr}})^2/(1+k_{\text{ibr}})^2 + 0.05$ has a gradient with respect to k_{ibr}:

$$\frac{dk_m}{dk_{\text{ibr}}} = \frac{-2(1-k_{\text{ibr}})}{(1+k_{\text{ibr}})^3} \leq 0.5 \quad \text{for all } k_{\text{ibr}} \in [0,1]$$

An estimation error of $|\Delta k_{\text{ibr}}| = 0.024$ pu produces a slope error of at most $0.5 \times 0.024 = 0.012$ in k_m. Since the stability margin $\eta = k_m - 2\epsilon_{\text{CT}} \geq 0.05 - 0.02 = 0.03$, the slope error (0.012) is **below the stability margin (0.03)**. The relay therefore cannot false-trip due to k_{ibr} estimation error alone.

Why the RMSE is small: The exponential filter with $\lambda = 0.10$ smooths out sample-to-sample noise while tracking trends in k_{ibr} within ≈ 10 samples (10 ms). The received phasor $\hat{I}_R$ from the remote relay over GOOSE is accurate to $< 0.5\%$ (IEC 61869 Class 1P CT accuracy). The dominant RMSE source is the filter's tracking lag during rapid k_{ibr} transitions (e.g., IBR FRT initiation), not measurement noise.

Stage-2 Gate: Only 1–2 Cycle Overhead

Why 1–2 Cycles is Acceptable

The Stage-2 gate adds processing time between the 87L criterion being satisfied and the trip command being issued:

- **LM estimation convergence**: 3–5 iterations of the LM algorithm, each requiring one Jacobian evaluation. At 1 kHz sampling: ≤ 5 ms (one quarter cycle).
- **Confidence score computation**: analytical formula, < 1 ms.
- **Gate decision**: single comparison $\gamma \geq 0.85$, < 0.1 ms.

Total Stage-2 overhead: < 6 ms per decision cycle. A power cycle at 50 Hz is 20 ms; one cycle = 20 ms, two cycles = 40 ms. The "1–2 cycle" statement means the **worst-case overhead is 40 ms**.

Why 40 ms overhead is acceptable:

1. Zone 1 instantaneous protection operates in 20 ms. Adding 40 ms overhead gives 60 ms total — still within the IEC 60255-181 *fast-trip* window of 80 ms.

2. The average trip time *decreases* by 1 ms (22 ms → 21 ms) because the adaptive slope lowers k_m at high k_{ibr}, making the operate criterion easier to meet earlier.
3. The 1–2 cycle overhead only applies to the Stage-2 check. The 87L criterion check (Stage-1) continues at 1 kHz with no delay.

IEC 60255 reference: IEC 60255-181 specifies a protection operating time limit of 80 ms for transmission line protection. SAMBP-87L with Stage-2 overhead: average 21 ms, maximum $\leq$ 60 ms — both well within compliance.

Robustness at k-ibr = 1.0: The Fully IBR Grid Case

The k-ibr = 1.0 Test: Most Challenging Scenario

$k_{\text{ibr}} = 1.0$ means the total fault current at the remote end equals the IBR rated current: $I_B = 1.0 \times I_{\text{rated}}$. The grid-side current I_A is from the SG (or grid infeed). At $k_{\text{ibr}} = 1.0$, the minimum operate-to-bias ratio is:

$$\rho_{\min} = \frac{2(1 - 1.0)}{1 + 1.0} = 0$$

The OBR approaches zero in the worst-case phase angle (180° phase opposition). In practice, the fault angle is not exactly 180°, so $\rho_{\min}$ is small but non-zero (typically 0.05–0.15 pu for realistic fault impedances).

SAMBP-87L at $k_{\text{ibr}} = 1.0$:

- Adaptive slope: $k_m(1.0) = 0.5 \times 0 + 0.05 = k_{\text{floor}} = 0.05$.
- Stability condition: $\epsilon_{\text{CT}} \leq (0.05 - I_c/I_{\text{bias}})/2$. For $I_c/I_{\text{bias}} \approx 0.01$: $\epsilon_{\text{CT}} \leq 0.02 = 2\%$. This requires Class 1P CTs (1% error limit) — a design requirement for lines with $k_{\text{ibr}} > 0.5$.
- The TDCS-87L provides the primary trip path at $k_{\text{ibr}} = 1.0$ for symmetric 3PH faults, exploiting the travelling-wave transient rather than the differential current magnitude.
- For asymmetric faults: 87LN activates with $\rho_2 = 2$ (independent of k_{ibr}).

Result: All 6 scenarios at $k_{\text{ibr}} = 1.0$ (3 fault types $\times$ 2 locations) are correctly tripped. TPR(at $k_{\text{ibr}} = 1.0$) = 1.000.

Monte Carlo Statistical Validation (TR-26)

The 110-scenario deterministic test demonstrates *correctness*. TR-26 demonstrates *statistical robustness* under parametric uncertainty:

Group	N	Correct	Rate	CI lo	CI hi	Target
A — 3PH internal (TPR)	5000	5000	1.000	0.9992	1.000	≥ 0.990
B — SLG internal (TPR)	5000	5000	1.000	0.9992	1.000	≥ 0.990
C — External (security)	5000	5000	1.000	0.9992	1.000	≥ 0.995

What the MC parameters varied: $k_{\text{ibr}} \in U[0.06, 0.20]$ pu; frequency $f \in U[47, 53]$ Hz; 3rd harmonic injection $I_{h3} \in U[0, 0.10]$ pu; CT ratio error $\epsilon_{\text{CT}} \in U[0, 0.005]$; pre-fault baseline drift $\hat{\mu}_{\text{pre}} \in U[0.002, 0.010]$ pu.

Key finding: All 15 000 Monte Carlo trials (3 groups $\times$ 5 000) are correctly classified. The Wilson 95% CI lower bound of 0.9992 means: even accounting for statistical uncertainty at $N = 5\,000$, the true performance rate is almost certainly above 99.9%.

Best Figure to Show at the TSA Meeting

Recommended Figure: TR-26 Figure 3 — Wilson CI Bar Chart

What it shows: Three bars, one per Monte Carlo group (Group A: 3PH internal TPR; Group B: SLG internal TPR; Group C: External security rate). All three bars reach exactly 1.000. Each bar has a Wilson 95% confidence interval whisker showing the CI lower bound at 0.9992. Two horizontal dashed lines mark the performance targets (≥ 0.990 for TPR, ≥ 0.995 for security). All bars exceed both targets by a large margin.

Why it is the best figure:

1. It proves *statistical rigour*: not just 110 scenarios correct, but 15 000 Monte Carlo trials correct with a CI lower bound of 0.9992.
2. It simultaneously validates TPR (sensitivity) and security (FPR = 0) in one compact figure.
3. The CI whiskers prove the result is not a lucky coincidence — even with the penalty of 5 000-trial statistics, the lower bound stays above 0.999.
4. Examiners familiar with statistical testing will immediately recognise the Wilson CI as a gold-standard approach.

Where to find it: TR-26/2026, page 4, Figure 3. Source: SAMBP test suite, `run_ibr_monte_carlo.py`, NumPy seed 2026.

Secondary figure: TR-26 Figure 1 (detection rate by k_{ibr} bin for Group A). This shows that every k_{ibr} bin from 0.06 to 0.20 pu achieves TPR = 1.000, with yellow bars (TDCS-87L only) and green bars (87L primary) clearly showing which element fires in which k_{ibr} regime.

Whiteboard supplement: Draw a single axis from $k_{\text{ibr}} = 0$ to 1.0. Mark three zones: (1) $k_{\text{ibr}} < 0.12$: TDCS dominant; (2) $0.12 < k_{\text{ibr}} < 0.54$: 87L primary with adaptive slope; (3) $k_{\text{ibr}} > 0.54$: 87LN + TDCS back-stop. Label the fixed relay's failure zone ($k_{\text{ibr}} > 0.54$). This makes the collaborative multi-element design immediately visual.

Cross-Thesis Connections

How the Results Connect Across the Thesis

These results validate upstream contributions:

- **Ch 2 (SAMBP Foundations)**: TPR = 1.000 and FPR = 0.000 empirically validates the Lyapunov stability proof and the security proposition of the bounded update law. The theoretical guarantees translate to zero misoperations in simulation.
- **TR-26 Group B element breakdown**: 87L detects 57.2% of SLG trials, 87LN detects 0.4%, 87LN0 detects 80.8%, TDCS detects 19.2%. This confirms the design intent that each element covers a distinct fault type regime.

These results enable downstream work:

- **Ch 6 (SAMBP-87T):** The TPR/FPR framework and Stage-2 gate timing analysis from Ch 3 are directly reused in Ch 6 for the transformer differential validation. The same test metric definitions allow cross-chapter comparison.
- **Ch 7 (WAPC):** The 87L TPR = 1.000 result is used by the WAPC ILP solver as a constraint: since 87L never misses an internal fault, it can serve as the "guaranteed back-stop" for all scenarios in the ILP formulation without adding a miss-probability penalty.
- **TR-35 (Coordination)**: The zero-gap coverage proven here is the 87L input to the full protection coordination matrix in TR-35, which demonstrates zero coverage gaps across all (k_{ibr}, α) combinations.

Comparing across SAMBP protection functions:

Function	TPR	FPR	Slide	Chapter
SAMBP-87L (line differential)	1.000	0.000	8	4
SAMBP-87T (transformer diff.)	0.991	0.009	10	6
SAMBP-SyncOC (generator OC)	1.000	0.000*	9	5
SAMBP-21 (adaptive distance)	0.960	0.040	11	7

* Spurious pickups reduced by 58%; FPR here refers to maloperation rate.

87L achieves the highest TPR/FPR combination because the 87L principle is inherently the most selective (current differential, not impedance or overcurrent).

Key Numbers You Must Know

Numbers to Have Ready for Any Examiner Question

Number	Value	Meaning
TPR (SAMBP-87L)	1.000	66/66 fault scenarios tripped
FPR (SAMBP-87L)	0.000	0/44 non-fault scenarios falsely tripped
TPR (fixed relay)	0.783	52/66 faults tripped (14 missed at high k_{ibr})
FPR (fixed relay)	0.148	6.5/44 non-faults falsely tripped
k_{ibr} RMSE	0.024 pu	LM estimator error < 2.4% of rated current
Trip time gain	−1 ms	22 ms → 21 ms (adaptive slope trips earlier)
Stage-2 overhead	1–2 cycles	< 40 ms worst case; avg < 6 ms
MC trials (TR-26)	15 000	3 groups × 5 000 trials
CI lower bound	0.9992	95% Wilson CI; target ≥ 0.990
k_m at $k_{\text{ibr}} = 1$	0.05	Floor value; requires Class 1P CT
k_{ibr} threshold	0.54	Fixed relay ($k_m = 0.30$) fails above this
TDCS detection	$k_{\text{det}} \geq 0.040$ pu	From TR-26 robustness analysis
External margin	15.5%	TDCS security margin at worst CT condition

Speaker Script (Timed to 3 Minutes)

What to Say at the TSA Meeting

[0:00–0:30] Set the context:
"This slide presents the validated simulation results for SAMBP-87L across a systematic 110-scenario test matrix. The test covers the full IBR operating envelope — k_{ibr} from 0.0 to 1.0 in steps of 0.1, three fault types (AG, BCG, and three-phase), and CT error stress-tested up to 5%. The 4-bus test network represents a realistic sub-transmission configuration with mixed synchronous and inverter sources."

[0:30–1:15] Present the numbers:
"The results are in the table. The fixed relay achieves a true positive rate of 0.783 — it misses 21.7% of faults, all at k_{ibr} above 0.54 where the operate-to-bias ratio falls below the fixed slope. The false positive rate is 0.148 — 14.8% of non-fault conditions (heavy load and CT error scenarios) produce spurious trips. SAMBP-87L achieves TPR = 1.000 and FPR = 0.000. Perfect discrimination. The average trip time actually decreases by 1 ms because the lower adaptive slope means the operate criterion is satisfied earlier in the fault transient. The k_{ibr} RMSE is 0.024 pu — the LM estimator tracks the IBR level to within 2.4% of rated current, which translates to a slope error of at most 0.012 — well within the 0.03 stability margin."

[1:15–2:00] Defend the statistical basis:
"The 110-scenario test is a point estimate. The real statistical validation is in TR-26 — an IBR-parametric Monte Carlo study with 5 000 trials per group. Both TPR groups (3PH and SLG internal faults) and the external security group achieve rate = 1.000. The Wilson 95% confidence interval lower bound is 0.9992 — far above the IEC 60255 target of 0.990. This means: with 95% statistical confidence, the true performance rate exceeds 99.92%. The best figure to reference here is TR-26 Figure 3 — the three-bar Wilson CI chart where all three bars hit 1.000 and the CI whiskers are visible but tight."

[2:00–2:40] Address robustness to k-ibr = 1.0:
"The slide specifically calls out robustness up to $k_{\text{ibr}} = 1.0$ — the fully IBR grid. At this extreme, the minimum operate-to-bias ratio approaches zero. The fixed relay cannot trip. SAMBP-87L uses the adaptive slope at its floor value of 0.05, and falls back on TDCS-87L for symmetric faults and 87LN for asymmetric faults. All 6 scenarios at $k_{\text{ibr}} = 1.0$ are correctly tripped. The Stage-2 gate adds at most 1–2 power cycles of overhead to the trip decision, for a worst-case total trip time of 60 ms — still within the IEC 60255-181 fast-trip window of 80 ms."

[2:40–3:00] Close with the key message:
"The bottom line: SAMBP-87L achieves perfect discrimination — TPR = 1.000 and FPR = 0.000 — across the full IBR operating envelope from a conventional network to a fully inverter-based grid, statistically confirmed by 15 000 Monte Carlo trials. The fixed-slope relay, by contrast, produces a 14.8% false-trip rate and misses more than one in five faults at high IBR penetration. This is the core quantitative case for adaptive line differential protection."

Anticipated Questions and Answers

Q1: How can average trip time decrease when you have added Stage-2 overhead

Answer: Two effects work in opposite directions. The Stage-2 gate adds up to 6 ms overhead. But the adaptive slope reduces k_m at high k_{ibr}, making the 87L operate criterion easier to satisfy. For $k_{\mathrm{ibr}} = 0.6$: fixed relay must wait for ρ to exceed 0.30 (which only happens when the fault current builds up to a sufficient level), taking on average 22 ms. SAMBP-87L with $k_m = 0.18$ at $k_{\mathrm{ibr}} = 0.6$ operates as soon as $\rho > 0.18$ — typically 4–6 ms earlier. The net effect: Stage-2 overhead (+6 ms) is outweighed by the earlier criterion satisfaction (-7 ms) for the high-k_{ibr} scenarios that make up the majority of the test matrix. Hence average trip time decreases by 1 ms.

Q2: Why is k-ibr RMSE reported rather than maximum error

Answer: RMSE is reported because it directly enters the slope error bound: slope error $\leq (dk_m/dk_{\mathrm{ibr}}) \times \Delta k_{\mathrm{ibr}}$. The maximum slope gradient is 0.5, so RMSE of 0.024 pu gives a slope error bound of 0.012 — which is below the 0.03 stability margin. The maximum single-sample error could be higher (up to 3–4 sigma $\approx$ 0.07 pu), but the exponential filter with $\lambda = 0.10$ means any single-sample spike is attenuated by 90% after one update. The effective slope error from a single spike is $0.10 \times 0.5 \times 0.07 = 0.004$ — negligible. Reporting RMSE correctly characterises the filter's steady-state tracking performance.

Q3: FPR = 0.000 in 44 non-fault scenarios. Is 44 sufficient to claim zero FPR

Answer: The 44-scenario point estimate alone is not sufficient — as a point estimate, it only confirms no false trips in the tested conditions. The TR-26 Monte Carlo Group C (external security) addresses this directly: 5 000 independent external fault and through-load trials with no false trips. The Wilson CI lower bound of 0.9992 provides the statistical guarantee. Equivalently: the true FPR is less than $1 - 0.9992 = 0.0008$ with 95% confidence. Even accounting for scenarios not tested (different load shapes, exotic IBR control modes), the Stage-2 gate provides a structural guarantee: unless $\gamma \geq 0.85$ is satisfied, the trip is blocked. The gate is a hard veto, not a probabilistic one.

Q4: The test uses CT error up to 5 percent. IEC Class 1P is 1 percent. Why test to 5 percent

Answer: The 5% stress test goes beyond Class 1P to characterise the *degradation curve* — how does performance change as CT error increases beyond specification? The result is that SAMBP-87L maintains FPR = 0.000 all the way to 5% CT error because the Stage-2 gate detects CT saturation through the elevated Jacobian condition number κ_n. At 5% error, κ_n rises sufficiently to drop γ below 0.85, vetoing the trip. This is an important design safety property: even with a significantly degraded CT (5$\times$ worse than Class 1P), the relay does not false-trip. It does slow down — some legitimate fault trips may also be vetoed under these conditions, which is why Class 1P CTs are specified for $k_{\text{ibr}} > 0.5$ applications.

Q5: The Monte Carlo uses k-ibr only up to 0.20 pu in TR-26. The deterministic test goes to 1.0. Why the difference

Answer: The TR-26 Monte Carlo focuses on the low-k_{ibr} regime (0.06–0.20 pu) where the TDCS-87L is the primary discriminator (87L threshold not met). This is the most statistically challenging regime — faults are close to the detection threshold, making this where false negatives are most likely. The high-k_{ibr} regime (> 0.20 pu) is covered by the main 87L adaptive threshold, which has a larger margin and is less statistically sensitive. The deterministic 110-scenario test covers the full range to $k_{\text{ibr}} = 1.0$ because deterministic testing is needed to verify the extreme case, while Monte Carlo is most valuable where the margin is thin. Together, the two approaches provide complementary coverage: statistical rigour at the margin, deterministic verification at the extreme.

Chapter 3 Results One-Line Summary for Examiners

"SAMBP-87L: TPR = 1.000, FPR = 0.000 across 110 scenarios from $k_{\text{ibr}} = 0$ *to 1.0 (AG, BCG, ABC; CT error up to 5%). Fixed relay: TPR = 0.783 (14 missed faults at* $k_{\text{ibr}} > 0.54$*), FPR = 0.148 (7 false trips under CT error).* k_{ibr} *RMSE = 0.024 pu; Stage-2 overhead = 1–2 cycles. TR-26 Monte Carlo (15 000 trials): all groups at 1.000 with Wilson CI lo = 0.9992."*